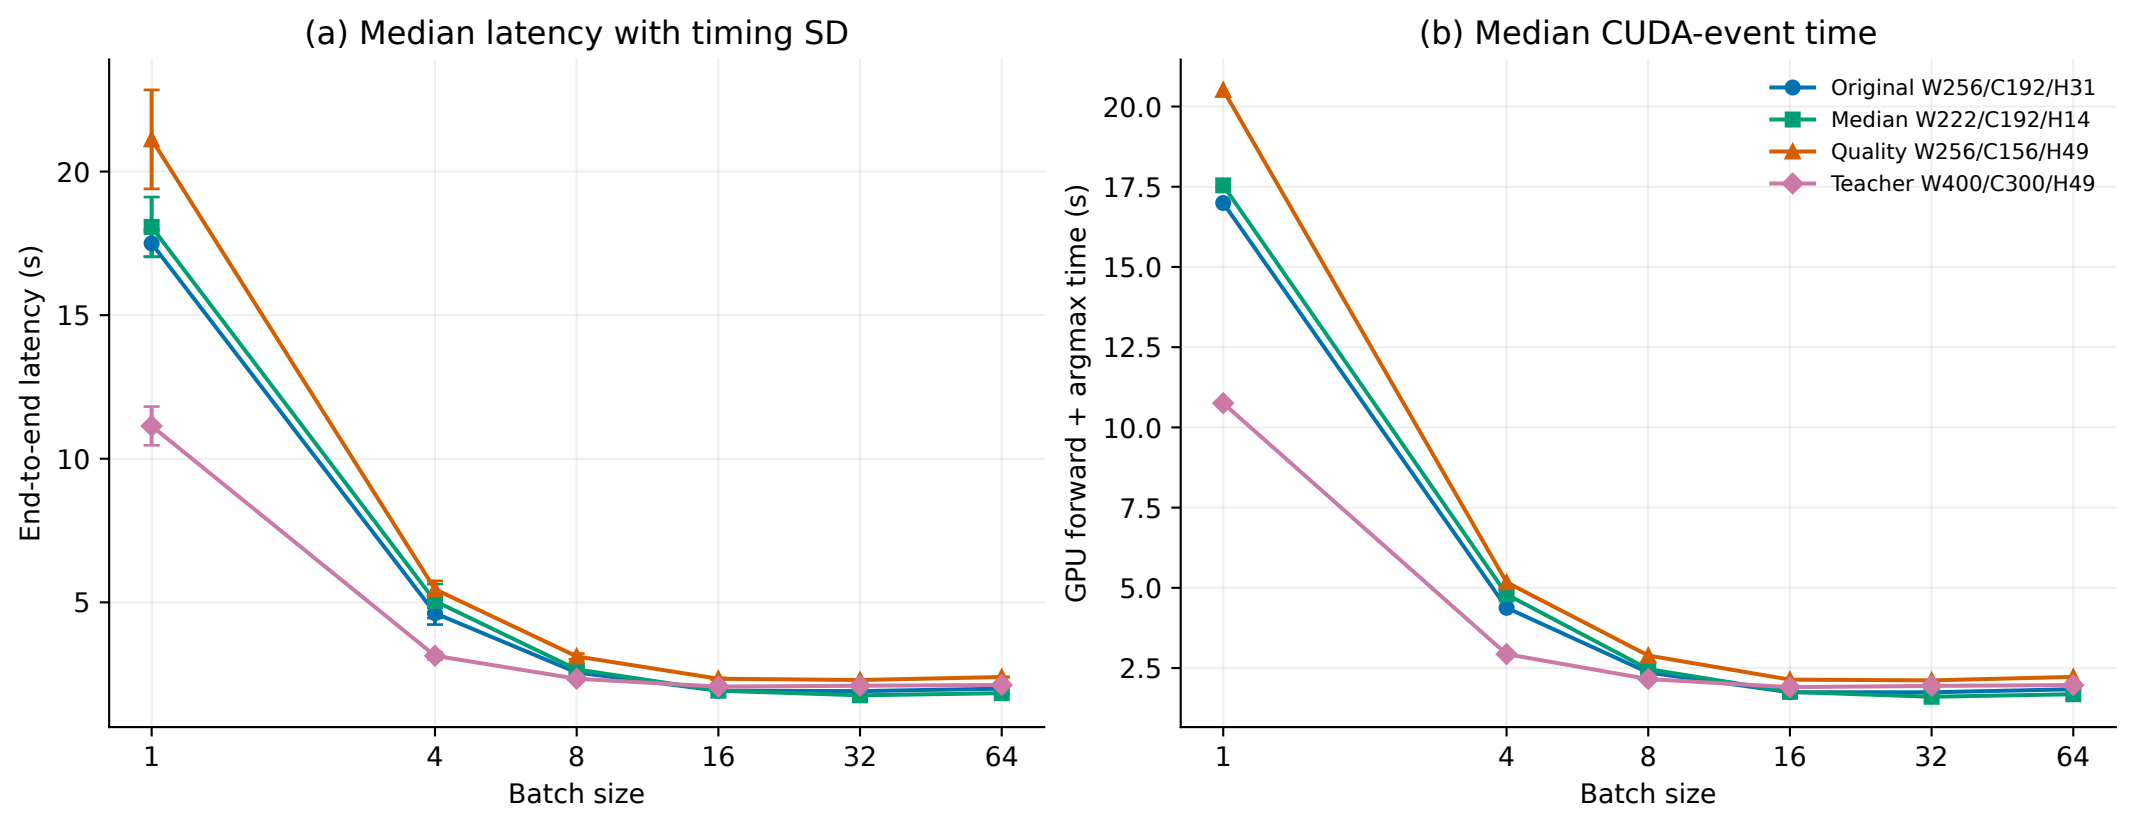
(a) Median latency with timing SD
End-to-end latency (s)
20
15
10
5
1
4
8
16
32
64
Batch size
(b) Median CUDA-event time
GPU forward + argmax time (s)
20.0
17.5
15.0
12.5
10.0
7.5
5.0
2.5
1
4
8
16
32
64
Batch size
Original W256/C192/H31
Median W222/C192/H14
Quality W256/C156/H49
Teacher W400/C300/H49