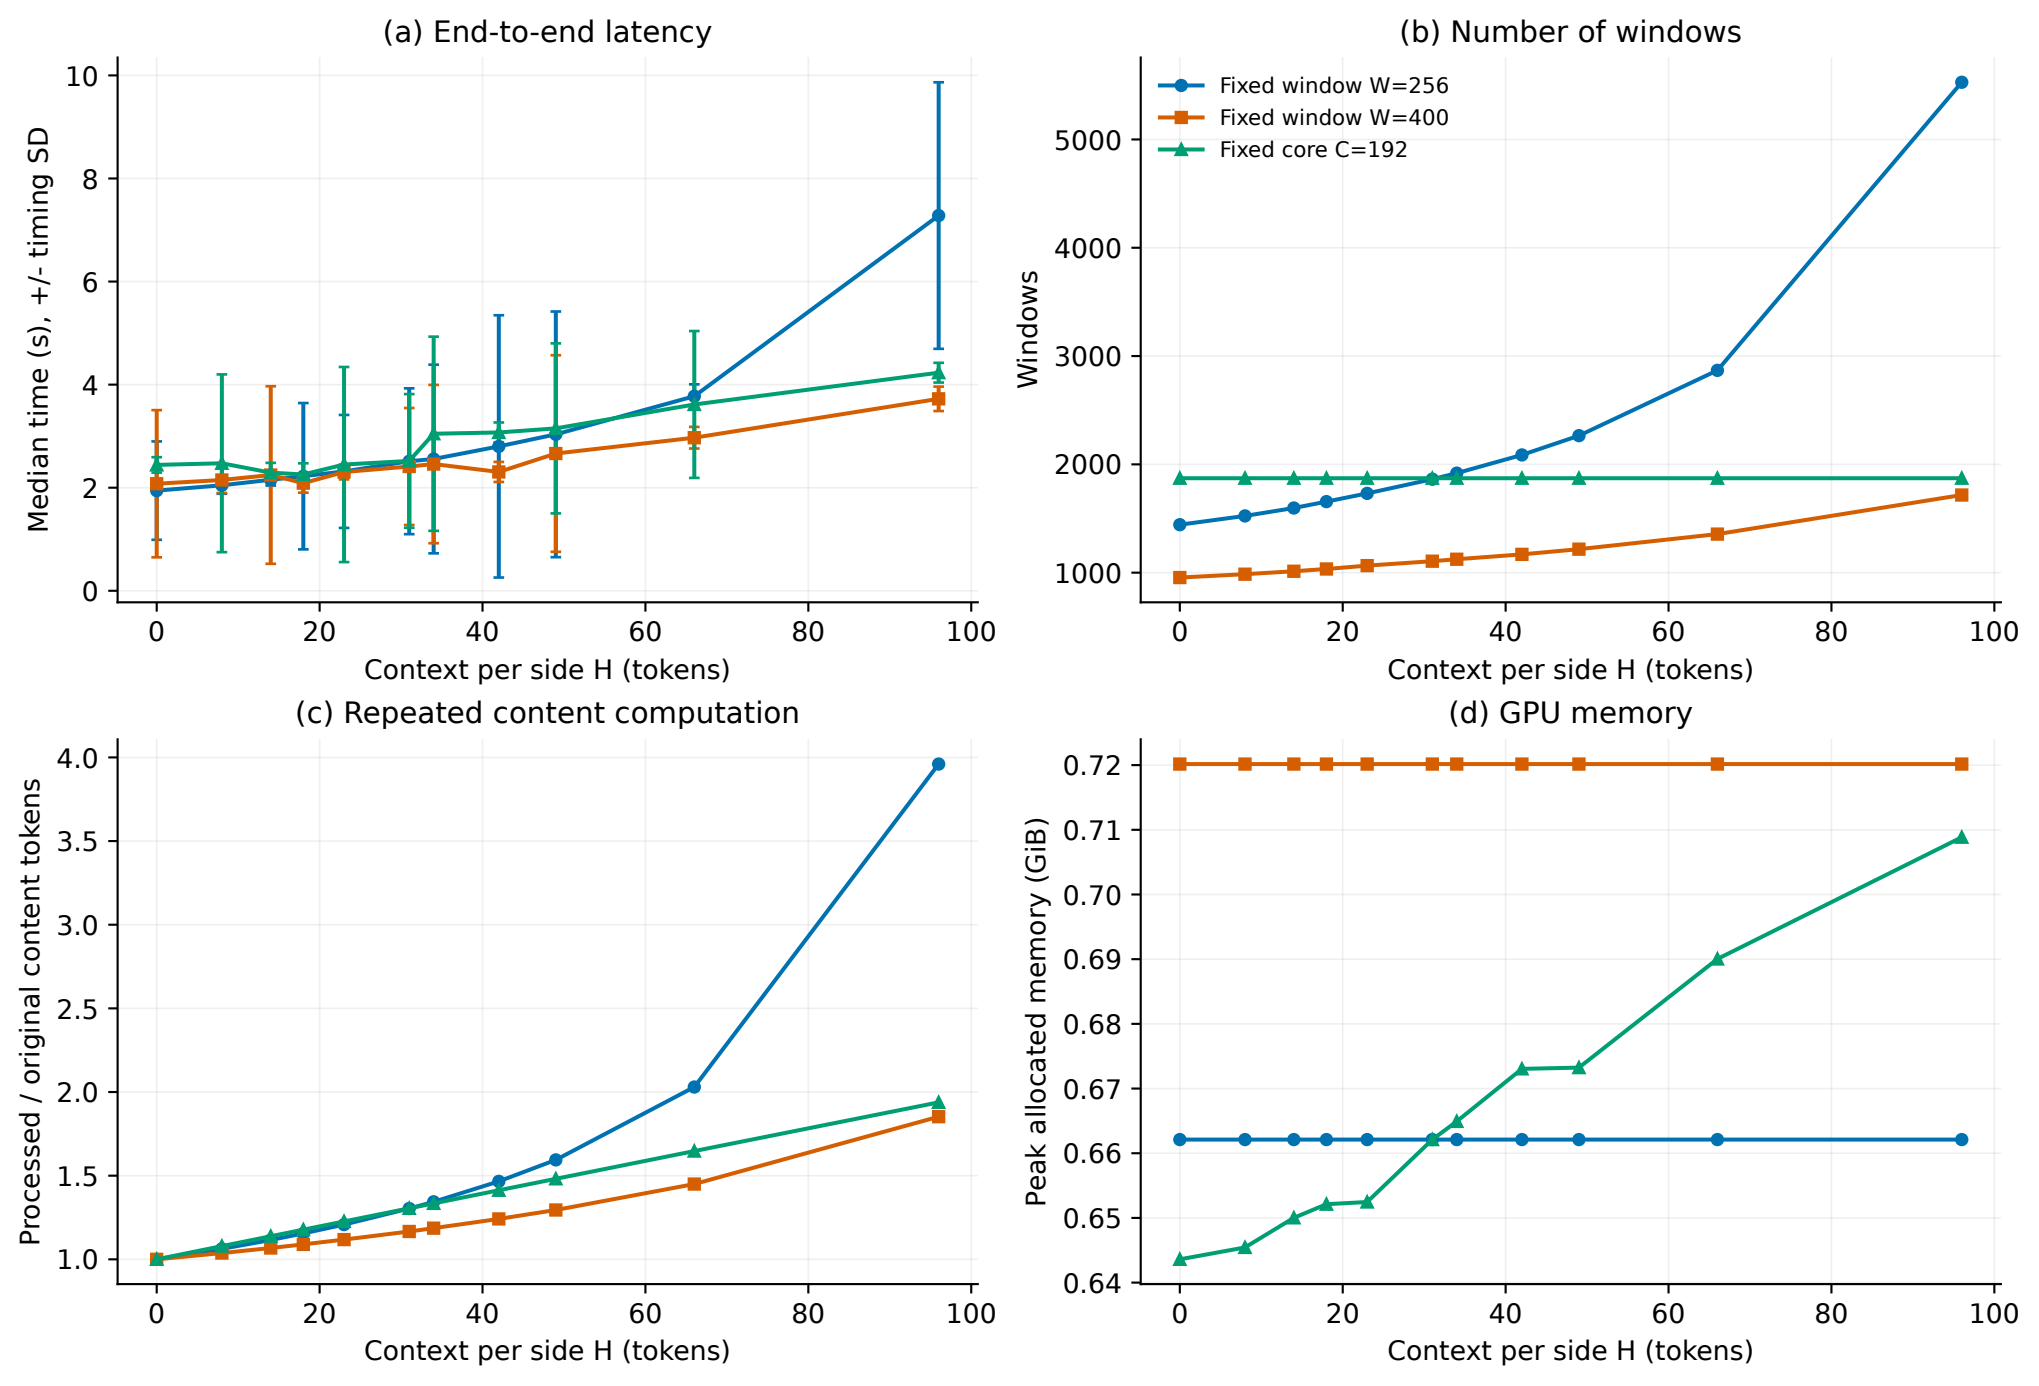
(a) End-to-end latency
Median time (s), +/- timing SD
Context per side H (tokens)
(b) Number of windows
Windows
Fixed window W=256
Fixed window W=400
Fixed core C=192
Context per side H (tokens)
(c) Repeated content computation
Processed / original content tokens
Context per side H (tokens)
(d) GPU memory
Peak allocated memory (GiB)
Context per side H (tokens)